

# Jacqueline Harding

[hardingi@stanford.edu](mailto:hardingi@stanford.edu) | [jacquelineharding.github.io](https://jacquelineharding.github.io)

AOS  
AOC

Philosophy of AI, including AI Ethics  
Logic and Formal Methods, Philosophy of Mind, Philosophy of Science

## EDUCATION

### Stanford University

PhD Philosophy and Symbolic Systems (expected June 2027)

### Institute for Logic, Language and Computation (ILLC), University of Amsterdam

MSc Logic and Computation (2019), *Cum Laude* (GPA: 9.1/10)

### Trinity College, University of Cambridge

MPhil Philosophy (2017), *Distinction* (ranked joint 1<sup>st</sup> in year)

BA Philosophy (2016), *First Class with Distinction* ("Starred First"; ranked 1<sup>st</sup> in year)

## NON-ACADEMIC RESEARCH EXPERIENCE

### Anthropic

Core Contributor, AI Moral Patienthood Project (December 2023–April 2024)

### Center for AI Safety (CAIS)

Research Fellow (January–August 2023)

## PUBLICATIONS (\*equal contribution)

### Journal Articles

A Communication-First Account of Explanation  
Jacqueline Harding, Tobias Gerstenberg, Thomas Icard  
*Noûs*, 2026

How Causal Abstraction Underpins Computational Explanation  
Atticus Geiger, Jacqueline Harding, Thomas Icard  
*Philosophy and the Mind Sciences*, 2026

What is AI Safety? What do we want it to be?  
Jacqueline Harding\*, Cameron Domenico Kirk-Giannini\*  
*Philosophical Studies*, 2025

What is it for a Machine Learning Model to have a Capability?  
Jacqueline Harding, Nathaniel Sharadin  
*The British Journal for the Philosophy of Science*, 2024

Operationalising Representation in Natural Language Processing  
Jacqueline Harding  
*The British Journal for the Philosophy of Science*, 2023

AI Language Models Cannot Replace Human Research Participants  
Jacqueline Harding\*, William D'Alessandro\*, N.G. Laskowski\*, Robert Long\*  
*AI and Society*, 2023

Proxy Selection in Transitive Proxy Voting  
Jacqueline Harding  
*Social Choice and Welfare*, 2022

Everettian Quantum Mechanics and the Metaphysics of Modality  
Jacqueline Harding  
*The British Journal for the Philosophy of Science*, 2021 (first published online 2018)

### Conference Proceedings

Toward a Formal Pragmatics of Explanation  
Jacqueline Harding, Tobias Gerstenberg, Thomas Icard  
*Proceedings of the Annual Meeting of the Cognitive Science Society*, 2025

Do As I Explain: Explanations Communicate Optimal Interventions  
Lara Kirfel, Jacqueline Harding, Jeong Yeon Shin, Cindy Xin, Thomas Icard, Tobias Gerstenberg  
*Proceedings of the Annual Meeting of the Cognitive Science Society*, 2024

Under the Hood: Using Diagnostic Classifiers to Investigate and Improve how Language Models Track Agreement Information  
Mario Giulianelli\*, Jacqueline Harding\*, Florian Mohnert\*, Dieuwke Hupkes, Willem Zuidema  
*Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, 2018

### Preprints

Taking AI Welfare Seriously  
Robert Long, Jeff Sebo, Patrick Butlin, Kathleen Finlinson, Kyle Fish, Jacqueline Harding, Jacob Pfau, Toni Sims, Jonathan Birch, David Chalmers  
*arXiv preprint*, 2024

### PUBLIC WRITING (\*equal contribution)

Toward a More Expansive Perspective on AI Safety (Blog Post and Workshop Report)  
Jacqueline Harding\*, César Valenzuela\*  
*HAI and McCoy Family Center for Ethics in Society, Stanford University*, July 2024

AI's future worries us. So does AI's present. (Opinion Piece)  
Jacqueline Harding\*, Cameron Domenico Kirk-Giannini\*  
*The Boston Globe*, July 2023

### PRESENTATIONS (+invited, \*reviewed)

**Comments on Laurie Paul's 'Knowledge in Animals and Machines' and Raphaël Millière's 'Tacit Knowledge of Rules in Large Language Models'**

+Symposium on 'Knowledge in Large Language Models', Central APA (February 2026)

**Comments on Cameron Buckner's 'From Deep Learning to Rational Machines'**

+Book Symposium, Pacific APA (April 2025)

### Goal-Directedness in AI Systems

+Session on AI Safety, Central APA (February 2026)  
+PAINT online lecture series, Northeastern University (March 2025)  
\*MINT/Yale Workshop on Normative Philosophy of Computing, Yale Law School (September 2024)

## **A Communication-First Account of Explanation**

+CSLI Workshop, Stanford University (June 2024)

## **What is it for a Machine Learning Model to have a Capability?**

+Ability Graduate Seminar, NYU (September 2024)  
+AI Benchmarking Workshop, The University of Hong Kong (March 2024)  
+Machine Intelligence and Normative Theory (MINT) Lab, Australian National University (November 2023)

## **Representation in Natural Language Processing**

+Philosophy of Mind Reading Group, NYU (September 2024)  
\*NYU/Columbia Philosophy of Deep Learning Conference, NYU (March 2023)

## **How do Language Models Track Agreement Information?**

\*BlackboxNLP Workshop, Empirical Methods in Natural Language Processing (EMNLP) Conference (November 2018)

## **TEACHING EXPERIENCE** (\*primary instructor)

### **Stanford University**

\*PHIL 24H: Philosophy of Large Language Models (Autumn 2024)  
PHIL 80: Mind, Matter and Meaning (Autumn 2024)  
PHIL 151/251: Metalogic (Winter 2022)  
PHIL 150/250: Mathematical Logic (Autumn 2021)

### **University of Amsterdam**

Natural Language Processing 1, Artificial Intelligence MSc (Autumn 2018)  
Mathematical Proof Methods for Logic, Logic MSc (grader, Autumn 2018)

### **University of Cambridge**

1A Logic, Philosophy BA (full academic year, 2016-2017)

## **SELECTED AWARDS**

### **Stanford University**

Mellon Foundation Dissertation Fellowship (2026)  
(awarded to ~10 final year Stanford PhD students across the humanities)  
  
Centennial Teaching Assistant Award (2022)  
(awarded to ~50 Stanford TAs each year, based on faculty and student evaluations)  
  
Human-Centered Artificial Intelligence (HAI) Graduate Fellowship (2021)  
(awarded to ~12 Stanford PhD students whose work intersects with AI each year)  
  
Patrick Suppes Fellowship in Philosophy of Science (2020-)

### **University of Amsterdam**

Amsterdam Science Talent Scholarship (2017-2019)  
(merit-based funding for MSc awarded to ~5 graduate students across the sciences)

### **University of Cambridge**

#### **Faculty of Philosophy**

Matthew Buncombe Prize (2017) (shared)

(awarded to the best performing student on the Philosophy MPhil)

Craig Taylor Prize (2016)

(awarded to the best performing student on the Philosophy BA)

### **Trinity College**

Travelling Studentship (2017)

(part funding for MSc)

Hyam Studentship (2016-2017)

(funding for MPhil)

Pre-Research Scholarship (2016), Senior Scholarship (2015), Junior Scholarship (2014)

(awarded on the basis of examination results in each year of the BA)

## **SERVICE**

### **Departmental**

#### **Faculty of Philosophy, Stanford University**

Founder and Organiser, Philosophy of AI Reading Group, 2024

Co-Head, Minorities and Philosophy (MAP) Chapter, 2022-2023

Mental Health Representative, 2021-2022

#### **ILLC, University of Amsterdam**

Master of Logic Mentor (mental health role), 2018-2019

#### **Trinity College, Cambridge**

Welfare Officer (mental health role), 2015-2016

### **Conference Organising**

Rapporteur, Sociotechnical AI Safety Conference

Stanford University (November 2023)

### **Refereeing**

Ethics and Information Technology, ACM FAccT Conference, Mind and Language, Minds and Machines, Philosophical Studies, Synthese